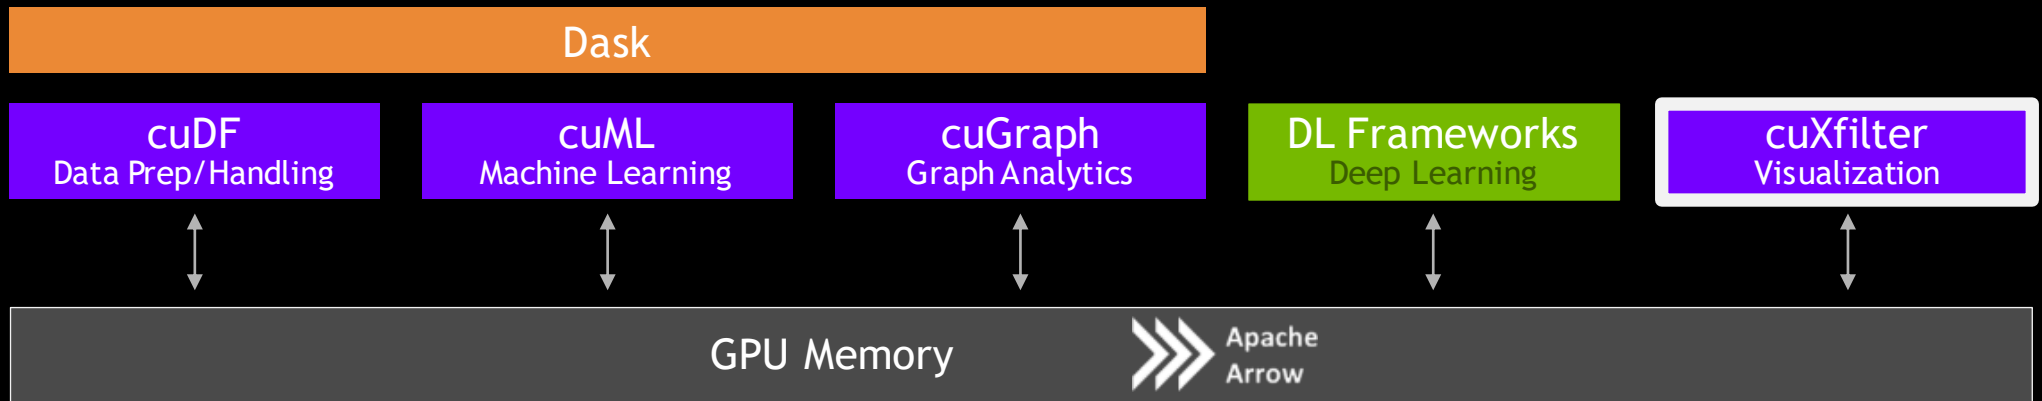




SECTION 2

01 - 07

RAPIDS PLATFORM



CUXFILTER

(coo-cross-filter)

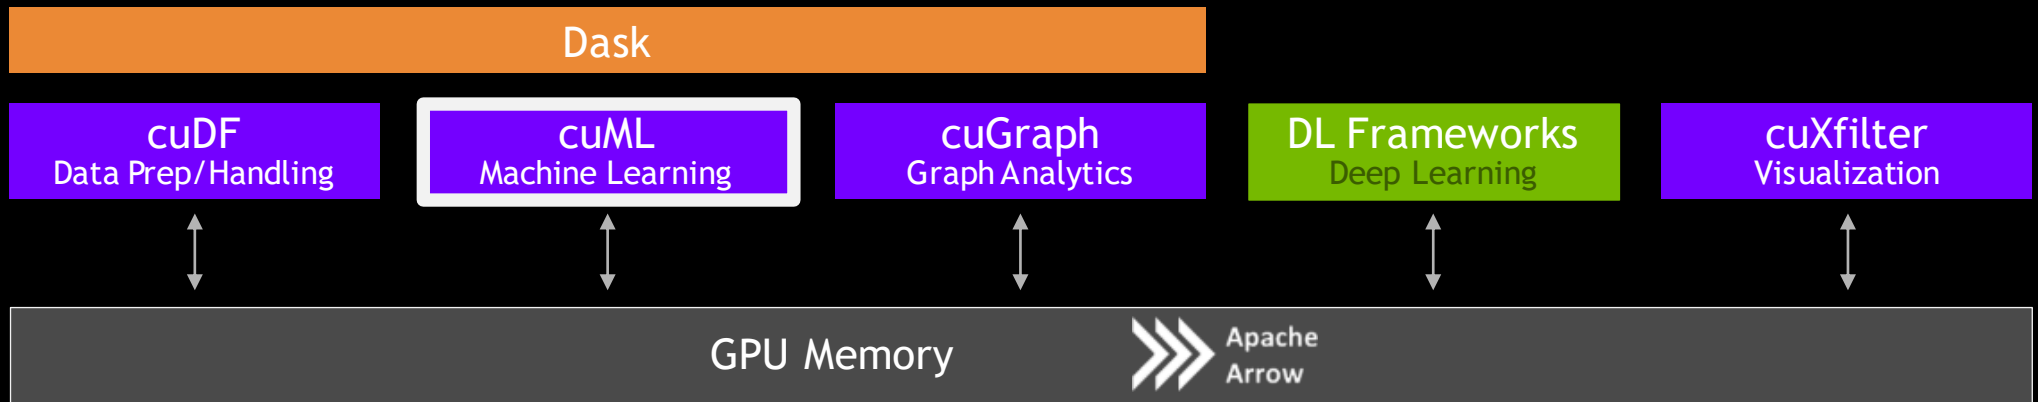
Efficient GPU-backed in-browser visualization engine

Built on a host of pyViz ecosystem tools

Bokeh and Datashader charts

Multiple charts with cross-filtering widgets

RAPIDS PLATFORM



CUML

Follows scikit-learn convention of model objects with `.fit` and `.predict/.transform`

Rapidly-growing subset of algorithms, driven by use cases

In today's workshop

- ▶ K-means (single- and multi-GPU)
- ▶ DBSCAN
- ▶ Logistic regression
- ▶ K-nearest neighbors
- ▶ XGBoost

K-MEANS

- ▶ Partitions data by iteratively moving cluster centers and reassigning datapoints based on which center is closest to the point
- ▶ Requires a known (or well-estimated) number of clusters
- ▶ Fast, simple, easy to understand
- ▶ Has an efficient multi-node, multi-GPU implementation

DBSCAN

- ▶ Uses spatial density to cluster—finds cluster points within at most *epsilon* distance of another point in the cluster
- ▶ **Can identify outliers**—not every point is “reachable” from a cluster core
- ▶ Does not require prior knowledge of the number of clusters, but does require an estimate for epsilon
- ▶ Can identify clusters of unusual shapes

K-NEAREST NEIGHBORS

- ▶ Enables rapid discovery of nearby, existing points to a new observation
- ▶ Algorithm fits a data structure for future use
- ▶ **Foundation of other algorithms** that require knowing which points are nearby a given point

LOGISTIC REGRESSION

- ▶ Regression with binary dependent variables
- ▶ Uses the logistic function to map $(-\infty, +\infty)$ domain scores to $(0, 1)$ range
- ▶ No simple coefficient formulas, as with ordinary least squares (OLS) linear regression
- ▶ Assumes independence of features

TRY NOTEBOOKS 01 - 07 NOW

docs.rapids.ai/api





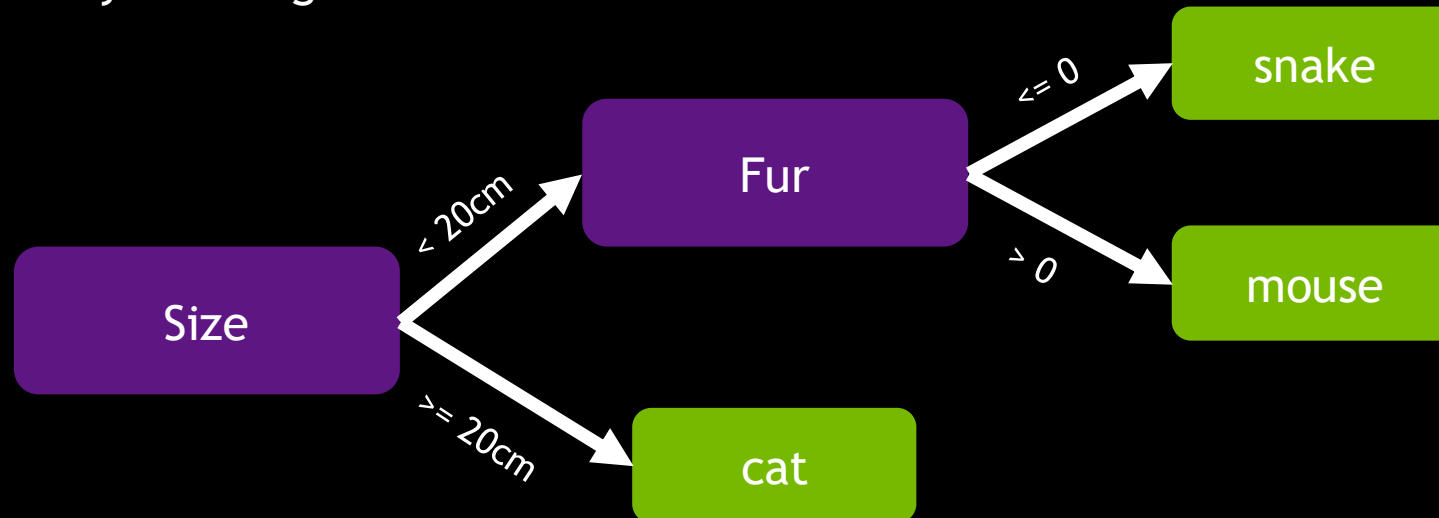
SECTION 2

08

XGBOOST

Popular and powerful general-purpose classification and regression algorithm for structured data—can identify and use complex patterns

Gradient-boosted decision trees—optimizing against an objective function while controlling tree complexity with regularization



TRY NOTEBOOK 08 NOW

docs.rapids.ai/api






SECTION 2

09

SINGLE-SOURCE SHORTEST PATH

Input: Graph (w/o negative-weight cycles), node_id

Output: dataframe with Vertex, Distance, Predecessor columns



`cg.sssp(G, 0)`

Vertex	Distance	Predecessor
0	0	-1
1	2	0
2	1	0
3	3	2

Note: using mapped integer IDs in exercise for memory and speed

TRY NOTEBOOK 09 NOW

docs.rapids.ai/api

